


Econometrics II - Problem 4

William Radaic Peron, Vinícius Marcial

August 28, 2020

In this problem, we'll be forecasting GDP in the short term and creating some models of GDP growth in the long run. This presents some challenges, namely those related to *ergodicity* and *stationarity*.

```
pplot <- ggplot(data = pib, aes(x = t, y = v)) + geom_line() + ggtitle("Time series plot") + theme_few(pplot)
```



As we have downloaded the *pure* quarterly data, it presents *seasonality* and an upwards tendency. This implies that the *time series will not be stationary*. Therefore, we need to employ methods that circumvent this issue and assure us that we can continue modelling the series as an ARMA(p,q).

Decomposing the time series

We will now assume that we can decompose the time series in three distinct elements in an additive model:

$$X_t = f_t + s_t + Y_t$$

, where f_t denotes the tendency of the ts, s_t denotes seasonality, Y_t is stochastic. We also assume that f_t, s_t are *deterministic*.

Trend

First, we'll construct a *parametric* model of the trend. Let's assume that f_t can be modelled by a linear form:

$$f_t = \gamma_0 + \gamma * t$$


```
linear_trend <- lm(v ~ t, data = pib )

summary(linear_trend)

##
## Call:
## lm(formula = v ~ t, data = pib)
##
## Residuals:
##      Min       1Q   Median       3Q      Max
## -34187 -11102  -1888   11332   32261
##
## Coefficients:
##              Estimate Std. Error t value Pr(>|t|)
## (Intercept) -1.225e+07  4.376e+05  -28.00   <2e-16 ***
## t              6.226e+01  2.179e+00   28.57   <2e-16 ***
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## Residual standard error: 15030 on 95 degrees of freedom
## Multiple R-squared:  0.8957, Adjusted R-squared:  0.8946
## F-statistic: 816.2 on 1 and 95 DF,  p-value: < 2.2e-16
ggplot(data = pib, aes(x = t, y = v)) + stat_smooth(method = "lm", se = F) + geom_line() + geom_point()

## `geom_smooth()` using formula 'y ~ x'
```

Linear trend, GDP





Another way to find f_t is via a *non-parametric* process. For this, we'll use an HP filter and a moving average.

```
pib_ts <- ts(pib$v)

hp_trend <- hpfilter(pib_ts, freq = 1600, type = "lambda")

plot(hp_trend)
```




Now, a moving average.

```
pib_ma <- ma(pib$v, order = 4)
```

```
autoplot(pib_ma, color = "blue") + geom_line(data = pib, aes(x = 1:length(pib$t), y = v), color = "red")
```

```
## Warning: Use of `pib$t` is discouraged. Use `t` instead.
```



Seasonality

We can now create a function for s_t . This will be done with dummies:

$$D_i = 1, i = t$$

$$D_i = 0 \text{ otherwise}$$

```
tri <- c(NA)
tri1 <- c(1,2,3,4)
i = 1
while (i < 25) {
  tri <- append(tri, tri1)
  i = i + 1
}
tri <- tri[-1]
tri <- c(tri, 1)
tri
```

```
## [1] 1 2 3 4 1 2 3 4 1 2 3 4 1 2 3 4 1 2 3 4 1 2 3 4 1 2 3 4 1 2 3 4 1 2
## [39] 3 4 1 2 3 4 1 2 3 4 1 2 3 4 1 2 3 4 1 2 3 4 1 2 3 4 1 2 3 4 1 2 3 4
## [77] 1 2 3 4 1 2 3 4 1 2 3 4 1 2 3 4 1 2 3 4 1
```

```
pib <- data.frame(pib, tri)
```

```
names(pib)[1] <- "t"
names(pib)[2] <- "v"
names(pib)[3] <- "tri"
```

```
dummies <- data.frame(matrix(NA, nrow = length(pib$t), ncol = 4))
```

```
for (j in 1:4) {
```

```
  dummies[j] <- as.numeric(pib$tri == j)
```

```
}
```

```
hp_fitted <- hp_trend[2]
```

```
hp_fitted <- hp_fitted$trend
```

```
detrend <- pib$v - hp_fitted
```

```
pib <- data.frame(pib, dummies, detrend)
```

```
names(pib) <- c("t", "v", "tri", "X1", "X2", "X3", "X4", "detrend")
```

```
head(pib)
```

```
##           t           v tri X1 X2 X3 X4  detrend
## 18 199601 170920.0    1  1  0  0  0 -7639.547
## 40 199602 176708.8    2  0  1  0  0 -2784.500
## 62 199603 189844.3    3  0  0  1  0  9422.081
## 84 199604 184112.9    4  0  0  0  1  2773.122
## 106 199701 176732.2    1  1  0  0  0 -5513.291
## 128 199702 185109.5    2  0  1  0  0  1969.025
```

```
dummy_lm <- lm(detrend ~ X2 + X3 + X4, data = pib)
```

```
summary(dummy_lm)
```

```
##
## Call:
## lm(formula = detrend ~ X2 + X3 + X4, data = pib)
##
## Residuals:
##      Min       1Q   Median       3Q      Max
## -13331.9  -2630.5    527.3   2556.1  10393.6
##
## Coefficients:
##              Estimate Std. Error t value Pr(>|t|)
## (Intercept)  -6237.1      901.3  -6.920 5.68e-10 ***
## X2             5975.6     1287.9   4.640 1.14e-05 ***
## X3            11595.9     1287.9   9.004 2.62e-14 ***
```

```
## X4          7636.9      1287.9    5.930 5.12e-08 ***
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## Residual standard error: 4507 on 93 degrees of freedom
## Multiple R-squared:  0.4752, Adjusted R-squared:  0.4583
## F-statistic: 28.07 on 3 and 93 DF,  p-value: 5.155e-13
```

Y_t


We'll now use the HP-filtered version of f_t and the dummy approach to s_t .

```
yt <- as.vector(pib$v) - (hp_fitted + dummy_lm$fitted.values)
```

```
mean(yt)
```

```
## [1] 3.300398e-12
```

```
autoplot(yt) + theme_few()
```



```
y <- data.frame(1:97, yt)
```

```
names(y) <- c("t", "yt")
```

```
y
```

```
##      t      yt
## 1    1 -1402.41369
```

```
## 2 2 -2523.01651
## 3 3 4063.33233
## 4 4 1373.37256
## 5 5 723.84187
## 6 6 2230.50849
## 7 7 3863.92989
## 8 8 4706.55898
## 9 9 -997.54878
## 10 10 1561.05940
## 11 11 1109.30140
## 12 12 -1474.84104
## 13 13 -3094.32198
## 14 14 -2711.06480
## 15 15 -3682.09728
## 16 16 -1089.22329
## 17 17 1013.09213
## 18 18 803.79363
## 19 19 1058.19431
## 20 20 3566.65440
## 21 21 3218.89028
## 22 22 867.89808
## 23 23 -2630.46624
## 24 24 -2296.64964
## 25 25 -845.87753
## 26 26 232.06901
## 27 27 369.65920
## 28 28 2345.81672
## 29 29 -1618.01826
## 30 30 -4443.53279
## 31 31 -4881.28103
## 32 32 -3306.87043
## 33 33 -1065.87009
## 34 34 953.31655
## 35 35 1160.97388
## 36 36 1647.85000
## 37 37 -842.42514
## 38 38 1941.99747
## 39 39 -3111.77504
## 40 40 -2844.19975
## 41 41 -1151.13141
## 42 42 -2698.84818
## 43 43 -2825.03717
## 44 44 -2142.66917
## 45 45 90.10818
## 46 46 1867.51533
## 47 47 527.30802
## 48 48 2927.33435
## 49 49 3770.18034
## 50 50 6606.99132
## 51 51 7355.57794
## 52 52 -5425.05905
## 53 53 -13331.94859
## 54 54 -10170.65008
## 55 55 -6764.16878
```



```
## 56 56 -2532.14379
## 57 57 -1481.53913
## 58 58 1105.53274
## 59 59 1349.60693
## 60 60 2722.57709
## 61 61 2556.14091
## 62 62 4745.84336
## 63 63 2502.03347
## 64 64 1608.28362
## 65 65 -547.79006
## 66 66 336.83230
## 67 67 3302.21456
## 68 68 3038.23959
## 69 69 2113.60387
## 70 70 7857.97391
## 71 71 8151.97684
## 72 72 7916.47809
## 73 73 10393.61558
## 74 74 5328.42838
## 75 75 5716.81118
## 76 76 7492.11790
## 77 77 6359.42698
## 78 78 -1634.22681
## 79 79 -5766.42171
## 80 80 -7518.98503
## 81 81 -6862.95259
## 82 82 -8927.03190
## 83 83 -10781.49950
## 84 84 -11823.38303
## 85 85 -3771.12417
## 86 86 -4660.66608
## 87 87 -4555.56332
## 88 88 -3823.95804
## 89 89 1304.36238
## 90 90 -930.39059
## 91 91 456.52095
## 92 92 34.62592
## 93 93 3158.11540
## 94 94 2259.66776
## 95 95 4010.86917
## 96 96 4898.07304
## 97 97 2311.58347
```

Identifying and estimating ARMA(p,q) for Y_t

We are now in a position to identify and estimate the best model for our time series Y_t .

Applying the function `auto.arima` from the package *forecast* to identify and estimate the model:

```
aa_model <- auto.arima(y$yt, num.cores = 24, max.d = 0, stepwise = F)

summary(aa_model)
```

```
## Series: y$yt
## ARIMA(2,0,0) with zero mean
```

```

##
## Coefficients:
##          ar1      ar2
##      0.9392 -0.2166
## s.e.  0.0984  0.0990
##
## sigma^2 estimated as 7502261:  log likelihood=-904.92
## AIC=1815.84   AICc=1816.09   BIC=1823.56
##
## Training set error measures:
##              ME      RMSE      MAE      MPE      MAPE      MASE      ACF1
## Training set 9.709132 2710.641 2027.254 55.27169 123.3816 0.9589014 0.005486639
print("t-values: ")

## [1] "t-values: "
aa_t <- matrix(NA, nrow = aa_model$arma[1] + aa_model$arma[2])

for (i in c(1:(aa_model$arma[1] + aa_model$arma[2]))) {
  aa_t[i] <- aa_model$coef[i]/sqrt(aa_model$var.coef[i,i])
}

aa_t <- data.frame(aa_t)

aa_t

##          aa_t
## 1  9.540429
## 2 -2.188490
aa_q <- Box.test(aa_model$residuals, lag = aa_model$arma[1] + aa_model$arma[2])
aa_q

##
## Box-Pierce test
##
## data:  aa_model$residuals
## X-squared = 0.0050343, df = 2, p-value = 0.9975
criteria <- matrix(NA, nrow = 1, ncol = 3)

aa_criteria <- data.frame("AR(2)*", aa_model$aic, aa_model$bic)

names(aa_criteria) <- c("Model", "AIC", "BIC")


aa_criteria

##      Model      AIC      BIC
## 1 AR(2)* 1815.835 1823.559
fac_e <- ggAcf(aa_model$residuals, type = "correlation", lag.max = 30, plot = T) + theme_few()
facp_e <- ggPacf(aa_model$residuals, type = "correlation", lag.max = 30, plot = T) + theme_few()


## Warning: Ignoring unknown parameters: type

```

fac_e




facp_e



```
mean(aa_model$residuals)
```

```
## [1] 9.709132
```

```
plot(aa_model$residuals)
```



```


facst <- ggAcf(y$yt, type = "correlation", lag.max = 30, plot = T) + theme_few()
fac1t <- ggAcf(y$yt, type = "correlation", lag.max = 5000, plot = T) + theme_few()

facpst <- ggPacf(y$yt, type = "correlation", lag.max = 30, plot = T) + theme_few()

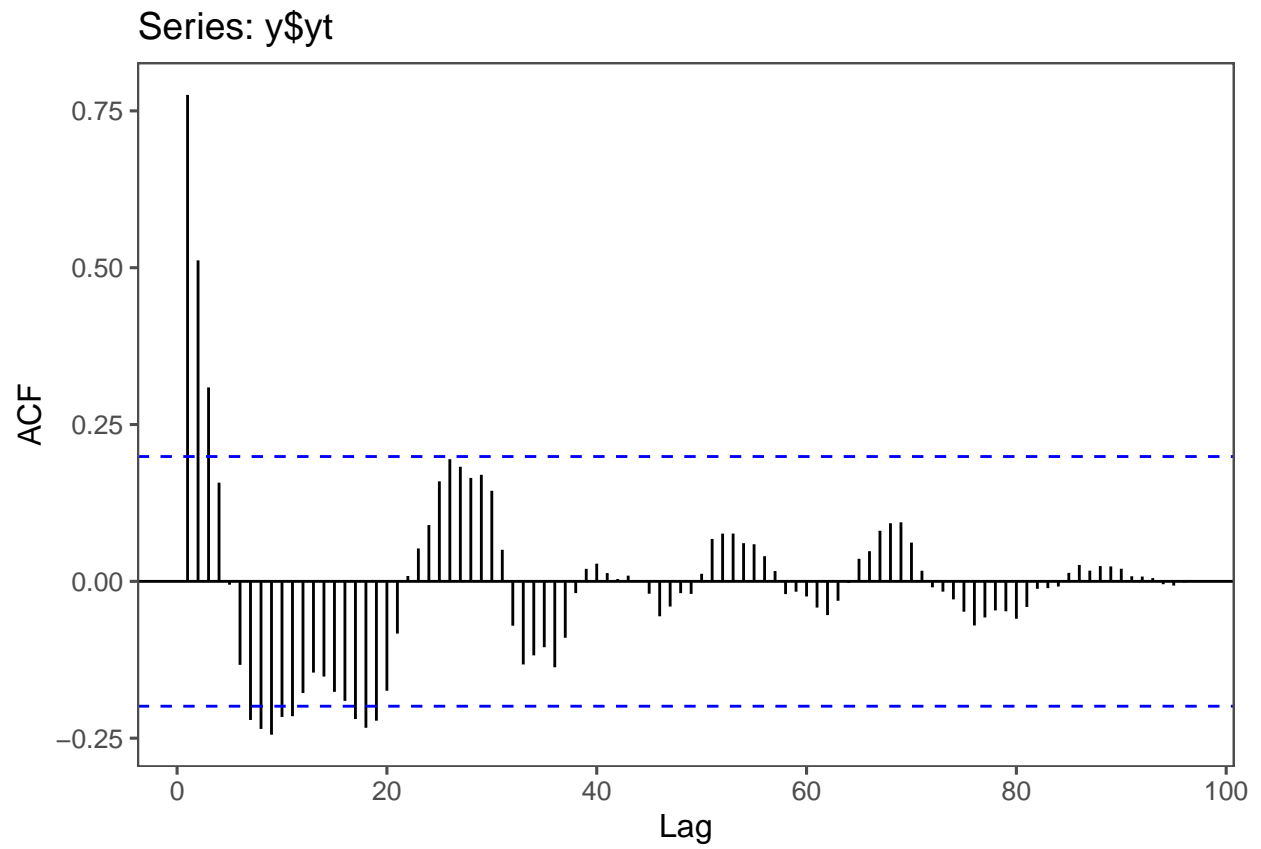
## Warning: Ignoring unknown parameters: type
facplt <- ggPacf(y$yt, type = "correlation", lag.max = 5000, plot = T) + theme_few()

## Warning: Ignoring unknown parameters: type
facst


```



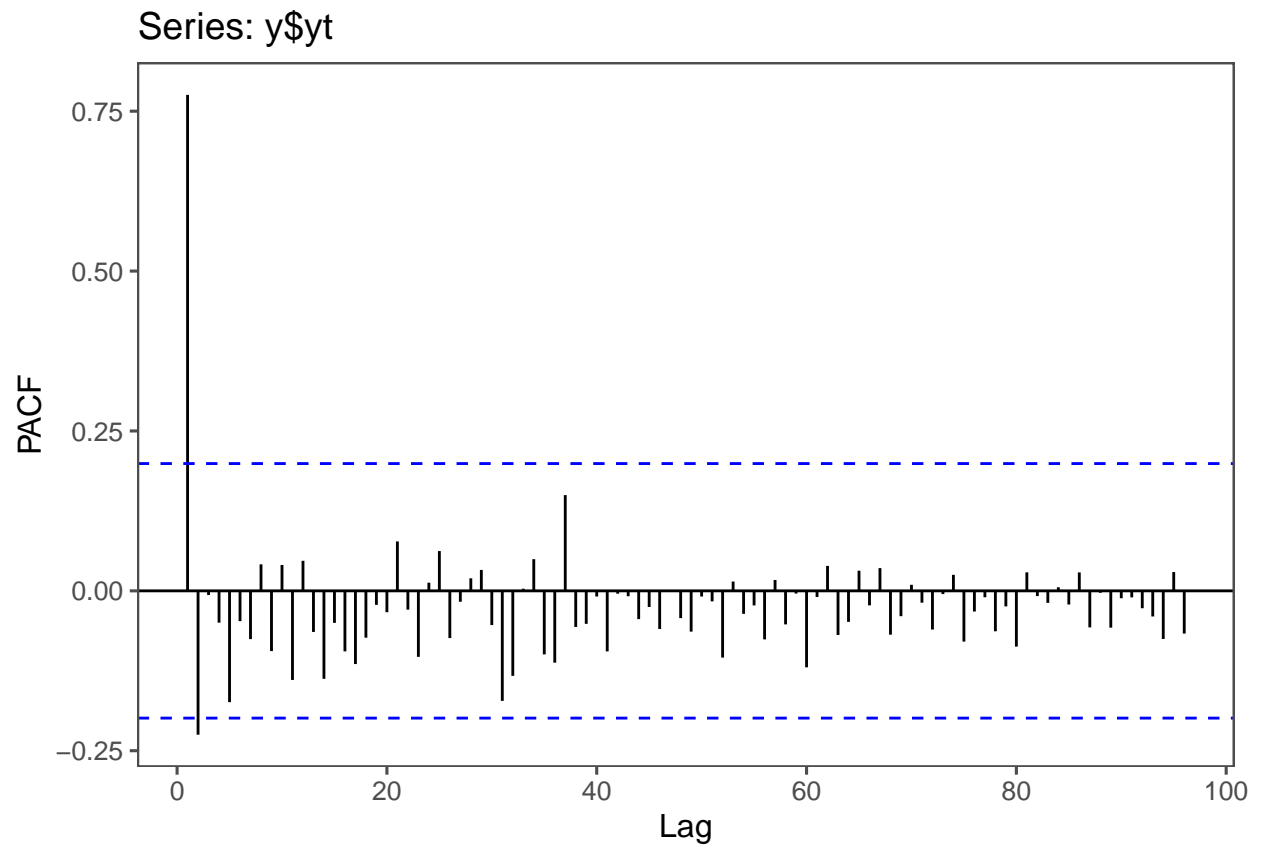
fac1t



facpst




facplt



The results of *auto.arima* imply that the best model is an ARMA(2,0) – i.e., an AR(2):

$$y_t = c + \phi_1 y_{t-1} + \phi_2 y_{t-2} + \varepsilon_t, \quad \varepsilon_t \sim wn(0, \sigma^2)$$

```
fc <- forecast(y$yt, model = aa_model, h = 4)
autoplot(fc) + theme_few()
```



Long term GDP growth

```
unemp <- read_excel("C:/Users/William/Downloads/tabela2176.xlsx")
```

```
## New names:
## * `` -> ...2
## * `` -> ...3
## * `` -> ...4
## * `` -> ...5
## * `` -> ...6
## * ...
```

```
unemp1 <- as.numeric(unemp[11,])
```

```
## Warning: NAs introduced by coercion
```

```
unemp2 <- unemp1[2:(length(unemp1)-2)]
```

```
unemp <- unemp2
```

```
df_unemp <- data.frame(1:length(unemp), unemp)
```

```
names(df_unemp) <- c("t", "r")
```

```
df_unemp
```

##	t	r
## 1	1	12.9
## 2	2	12.5
## 3	3	11.9
## 4	4	11.6
## 5	5	11.9
## 6	6	11.7
## 7	7	11.5
## 8	8	11.2
## 9	9	10.9
## 10	10	10.5
## 11	11	11.2
## 12	12	11.6
## 13	13	12.1
## 14	14	12.4
## 15	15	12.9
## 16	16	13.0
## 17	17	12.8
## 18	18	13.0
## 19	19	13.0
## 20	20	12.9
## 21	21	12.2
## 22	22	10.9
## 23	23	11.7
## 24	24	12.0
## 25	25	12.8
## 26	26	13.1
## 27	27	12.2
## 28	28	11.7
## 29	29	11.2
## 30	30	11.4
## 31	31	10.9
## 32	32	10.5
## 33	33	10.7
## 34	34	9.6
## 35	35	10.2
## 36	36	10.7
## 37	37	10.8
## 38	38	10.8
## 39	39	10.2
## 40	40	9.4
## 41	41	9.4
## 42	42	9.4
## 43	43	9.6
## 44	44	9.6
## 45	45	9.6
## 46	46	8.4
## 47	47	9.3
## 48	48	10.1
## 49	49	10.4
## 50	50	10.4
## 51	51	10.2
## 52	52	10.4
## 53	53	10.7


##	54	54	10.6
##	55	55	10.0
##	56	56	9.8
##	57	57	9.5
##	58	58	8.4
##	59	59	9.3
##	60	60	9.9
##	61	61	10.2
##	62	62	10.2
##	63	63	10.2
##	64	64	9.7
##	65	65	9.5
##	66	66	9.5
##	67	67	9.0
##	68	68	8.7
##	69	69	8.2
##	70	70	7.4
##	71	71	8.0
##	72	72	8.7
##	73	73	8.6
##	74	74	8.5
##	75	75	7.9
##	76	76	7.8
##	77	77	8.1
##	78	78	7.6
##	79	79	7.7
##	80	80	7.5
##	81	81	7.6
##	82	82	6.8
##	83	83	8.2
##	84	84	8.5
##	85	85	9.0
##	86	86	8.9
##	87	87	8.8
##	88	88	8.1
##	89	89	8.0
##	90	90	8.1
##	91	91	7.7
##	92	92	7.5
##	93	93	7.3
##	94	94	6.8
##	95	95	7.2
##	96	96	7.4
##	97	97	7.6
##	98	98	7.3
##	99	99	7.4
##	100	100	7.0
##	101	101	6.9
##	102	102	6.7
##	103	103	6.2
##	104	104	6.0
##	105	105	5.7
##	106	106	5.3
##	107	107	6.0

##	108	108	6.3
##	109	109	6.4
##	110	110	6.4
##	111	111	6.3
##	112	112	6.2
##	113	113	6.0
##	114	114	6.0
##	115	115	6.0
##	116	116	5.7
##	117	117	5.2
##	118	118	4.7
##	119	119	5.5
##	120	120	5.7
##	121	121	6.2
##	122	122	6.0
##	123	123	5.8
##	124	124	5.9
##	125	125	5.4
##	126	126	5.3
##	127	127	5.4
##	128	128	5.3
##	129	129	4.9
##	130	130	4.6
##	131	131	5.5
##	132	132	5.6
##	133	133	5.7
##	134	134	5.8
##	135	135	5.8
##	136	136	6.0
##	137	137	5.6
##	138	138	5.3
##	139	139	5.4
##	140	140	5.2
##	141	141	4.6
##	142	142	4.3
##	143	143	4.8
##	144	144	5.1
##	145	145	5.0
##	146	146	4.8
##	147	147	4.9
##	148	148	4.8
##	149	149	4.9
##	150	150	5.0
##	151	151	4.8
##	152	152	4.6
##	153	153	4.8
##	154	154	4.3
##	155	155	5.3
##	156	156	5.8
##	157	157	6.1
##	158	158	6.4
##	159	159	6.7
##	160	160	6.9
##	161	161	7.5

```
## 162 162 7.5
## 163 163 7.5
## 164 164 7.8
## 165 165 7.5
## 166 166 6.9
```

```
hp_unemp <- hpfilter(df_unemp$r, freq = 1600, type = "lambda")
```

```
plot(hp_unemp)
```



```
nairu <- hp_unemp$trend
```

```
nairu
```

```
##           [,1]
## [1,] 11.927416
## [2,] 11.924286
## [3,] 11.921765
## [4,] 11.920818
## [5,] 11.922402
## [6,] 11.927269
## [7,] 11.936158
## [8,] 11.949668
## [9,] 11.968123
## [10,] 11.991379
## [11,] 12.018626
## [12,] 12.048121
```

```
## [13,] 12.077607
## [14,] 12.104550
## [15,] 12.126430
## [16,] 12.140909
## [17,] 12.146136
## [18,] 12.140793
## [19,] 12.123974
## [20,] 12.095308
## [21,] 12.054973
## [22,] 12.003649
## [23,] 11.942106
## [24,] 11.870426
## [25,] 11.788538
## [26,] 11.696453
## [27,] 11.594813
## [28,] 11.485140
## [29,] 11.369331
## [30,] 11.249419
## [31,] 11.127330
## [32,] 11.005086
## [33,] 10.884566
## [34,] 10.767331
## [35,] 10.654831
## [36,] 10.547783
## [37,] 10.446622
## [38,] 10.351876
## [39,] 10.264295
## [40,] 10.184909
## [41,] 10.114709
## [42,] 10.054192
## [43,] 10.003413
## [44,] 9.962014
## [45,] 9.929387
## [46,] 9.904698
## [47,] 9.886905
## [48,] 9.874028
## [49,] 9.863719
## [50,] 9.853772
## [51,] 9.842314
## [52,] 9.827816
## [53,] 9.808972
## [54,] 9.784832
## [55,] 9.755005
## [56,] 9.719608
## [57,] 9.678911
## [58,] 9.633236
## [59,] 9.582792
## [60,] 9.527016
## [61,] 9.465171
## [62,] 9.396752
## [63,] 9.321711
## [64,] 9.240505
## [65,] 9.154139
## [66,] 9.063905
```

```
## [67,] 8.971311
## [68,] 8.878138
## [69,] 8.786185
## [70,] 8.697138
## [71,] 8.612320
## [72,] 8.532241
## [73,] 8.457027
## [74,] 8.386911
## [75,] 8.322216
## [76,] 8.263334
## [77,] 8.210393
## [78,] 8.163232
## [79,] 8.121622
## [80,] 8.084981
## [81,] 8.052462
## [82,] 8.022855
## [83,] 7.994666
## [84,] 7.965636
## [85,] 7.933637
## [86,] 7.896871
## [87,] 7.854210
## [88,] 7.805152
## [89,] 7.749784
## [90,] 7.688379
## [91,] 7.621368
## [92,] 7.549437
## [93,] 7.473321
## [94,] 7.393725
## [95,] 7.311247
## [96,] 7.226110
## [97,] 7.138472
## [98,] 7.048596
## [99,] 6.957036
## [100,] 6.864503
## [101,] 6.771982
## [102,] 6.680546
## [103,] 6.591347
## [104,] 6.505548
## [105,] 6.424068
## [106,] 6.347512
## [107,] 6.276030
## [108,] 6.209117
## [109,] 6.146098
## [110,] 6.086352
## [111,] 6.029420
## [112,] 5.975035
## [113,] 5.923103
## [114,] 5.873667
## [115,] 5.826822
## [116,] 5.782738
## [117,] 5.741696
## [118,] 5.703924
## [119,] 5.669312
## [120,] 5.637122
```



```
## [121,] 5.606511
## [122,] 5.576675
## [123,] 5.547180
## [124,] 5.517857
## [125,] 5.488697
## [126,] 5.459926
## [127,] 5.431719
## [128,] 5.404147
## [129,] 5.377264
## [130,] 5.351057
## [131,] 5.325217
## [132,] 5.298963
## [133,] 5.271624
## [134,] 5.242719
## [135,] 5.212032
## [136,] 5.179698
## [137,] 5.146216
## [138,] 5.112602
## [139,] 5.080152
## [140,] 5.050281
## [141,] 5.024603
## [142,] 5.004827
## [143,] 4.992394
## [144,] 4.988308
## [145,] 4.993449
## [146,] 5.008769
## [147,] 5.035225
## [148,] 5.073642
## [149,] 5.124759
## [150,] 5.189148
## [151,] 5.267236
## [152,] 5.359336
## [153,] 5.465466
## [154,] 5.585171
## [155,] 5.717579
## [156,] 5.861015
## [157,] 6.013544
## [158,] 6.173192
## [159,] 6.338039
## [160,] 6.506307
## [161,] 6.676444
## [162,] 6.847145
## [163,] 7.017617
## [164,] 7.187480
## [165,] 7.356650
## [166,] 7.525429
```

```
t6162 <- get_sidra(6612, variable = 9318, category = 90707, period = c("200202", "200203", "200204", "200205"))
```

```
## Considering all categories once 'classific' was set to 'all' (default)
```

```
View(t6162)
```

```
tax <- t6162[(t6162$`Setores e subsetores (Código)` == 90706),]
```

```

tax2 <- tax[,c(5,13)]

names(tax2)[1] <- "t"
names(tax2)[2] <- "r"

tax <- tax2

trimestra <- c(NA)

i <- 0
while (i<length(unemp)){

  media <- (unemp[i] + unemp[i+1] + unemp [i+2])/3
  trimestra <- append(trimestra, media)

  i<- i +3
}

nairu_3m <- trimestra

df <- data.frame(nairu_3m[-1], tax)

df

```

```

##      nairu_3m..1.      t      r
## 17      11.800000 200202 27105.64
## 39      11.466667 200203 27761.39
## 61      10.866667 200204 27645.46
## 83      12.033333 200301 26900.42
## 105     12.900000 200302 26877.27
## 127     12.966667 200303 27423.26
## 149     11.600000 200304 27890.17
## 171     12.633333 200401 27650.37
## 193     11.700000 200402 28727.94
## 215     10.933333 200403 29776.18
## 237     10.166667 200404 29890.40
## 259     10.766667 200501 28736.12
## 281      9.666667 200502 30275.09
## 303      9.533333 200503 30877.38
## 325      9.100000 200504 31098.96
## 347     10.300000 200601 30798.18
## 369     10.433333 200602 31768.30
## 391     10.133333 200603 32743.37
## 413      9.066667 200604 32347.13
## 435     10.100000 200701 32550.07
## 457      9.800000 200702 34105.20
## 479      9.066667 200703 34884.01
## 501      7.866667 200704 35832.30
## 523      8.600000 200801 35227.37
## 545      7.933333 200802 37054.11
## 567      7.600000 200803 38728.42
## 589      7.533333 200804 36680.04
## 611      8.800000 200901 34106.02
## 633      8.300000 200902 35886.78

```

```
## 655      7.766667 200903 38103.93
## 677      7.100000 200904 39175.00
## 699      7.433333 201001 38717.90
## 721      7.100000 201002 39961.02
## 743      6.300000 201003 41933.45
## 765      5.666667 201004 42540.35
## 787      6.366667 201101 41367.14
## 809      6.166667 201102 42736.91
## 831      5.900000 201103 43741.57
## 853      5.133333 201104 43929.18
## 875      5.966667 201201 42616.53
## 897      5.700000 201202 43727.51
## 919      5.333333 201203 45251.56
## 941      5.000000 201204 46492.43
## 963      5.700000 201301 43903.67
## 985      5.800000 201302 45932.85
## 1007     5.300000 201303 47239.43
## 1029     4.566667 201304 47669.43
## 1051     4.966667 201401 45634.31
## 1073     4.866667 201402 45631.55
## 1095     4.800000 201403 46892.87
## 1117     4.800000 201404 47982.22
## 1139     6.100000 201501 44453.25
## 1161     7.033333 201502 43600.46
## 1183     7.600000 201503 43585.04
## 1205      NA      201504 43359.17
```

```
names(df) <- c("NAIRU", "t", "tax")
```

```
growth_lm <- lm(NAIRU ~ tax, data = df)
```

```
summary(growth_lm)
```

```
##
## Call:
## lm(formula = NAIRU ~ tax, data = df)
##
## Residuals:
##      Min       1Q   Median       3Q      Max
## -1.1801 -0.3400 -0.0709  0.3167  1.6633
##
## Coefficients:
##              Estimate Std. Error t value Pr(>|t|)
## (Intercept)  2.110e+01  4.471e-01  47.19  <2e-16 ***
## tax         -3.479e-04  1.184e-05 -29.37  <2e-16 ***
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
## Residual standard error: 0.5991 on 52 degrees of freedom
## (1 observation deleted due to missingness)
## Multiple R-squared:  0.9431, Adjusted R-squared:  0.942
## F-statistic: 862.5 on 1 and 52 DF, p-value: < 2.2e-16
```

```
ggplot(data = df, aes(x = tax, y = NAIRU)) + stat_smooth(nethod = "lm", se = F) + theme_few()
```

```
## Warning: Ignoring unknown parameters: nethod
```

```
## `geom_smooth()` using method = 'loess' and formula 'y ~ x'
```

```
## Warning: Removed 1 rows containing non-finite values (stat_smooth).
```

